

2021

Cairo University

Faculty of computers and artificial intelligence



**AI322**

**Supervised Learning**

**Assignment\_3**

**CNN**

**Name: Gehad Mustafa Mansour Refae**

SUPERVISED LEARNING

STUD 20180080MNIST DATASET

## Default code:

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
from subprocess import check_output
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
import matplotlib.pyplot as plt
import matplotlib
import matplotlib.image as mpimg
import numpy as np
from numpy import random
import keras
from keras.utils import to_categorical
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten, Conv2D, MaxPooling2D
from keras.utils import np_utils
from keras.preprocessing.image import ImageDataGenerator
from keras.callbacks import ReduceLROnPlateau, Callback
from keras import regularizers
from keras.optimizers import Adam
import tensorflow as tf
from keras.optimizers import SGD
from sklearn.utils import shuffle
data_train = pd.read_csv('/content/mnist_train.csv')
data_test = pd.read_csv('/content/mnist_test.csv')
print("data_train" , data_train.shape)
print("data_test" , data_test.shape)
img_rows, img_cols = 28, 28
input_shape = (img_rows, img_cols, 1)
X = np.array(data_train.iloc[:, 1:])
y = to_categorical(np.array(data_train.iloc[:, 0]))
#Here we split validation data to optimize classifier during training
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,
random_state=13)
#Test data
X_test = np.array(data_test.iloc[:, 1:])
y_test = to_categorical(np.array(data_test.iloc[:, 0]))
X_train = np.array(X_train)
y_train = np.array(y_train)
X_test = np.array(X_test)
y_test = np.array(y_test)
X_train, y_train = shuffle(X_train, y_train)
X_test, y_test = shuffle(X_test, y_test)
X_train = X_train.reshape(X_train.shape[0], img_rows, img_cols, 1)
X_test = X_test.reshape(X_test.shape[0], img_rows, img_cols, 1)
X_val = X_val.reshape(X_val.shape[0], img_rows, img_cols, 1)
```

```

X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_val = X_val.astype('float32')
X_train /= 255
X_test /= 255
X_val /= 255
print(X_train)
print("X_train:{}\ny_train:{}\n\nX_val:{}\ny_val:{}\n\nX_test:{}\ny_test:{}".format
(X_train.shape, y_train.shape, X_val.shape, y_val.shape, X_test.shape,
y_test.shape))
batch_size = 32
num_classes = 10
epochs = 5
#input image dimensions
img_rows, img_cols = 28, 28

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
activation='relu',
kernel_initializer='he_normal',
input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
# opt = SGD(lr=0.001, momentum=0.9)
# opt = keras.optimizers.SGD(learning_rate=0.001)
opt = keras.optimizers.Adam(learning_rate=0.001)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=[
'accuracy'])

# model.compile(loss = 'categorical_crossentropy', optimizer = 'adam',
metrics = ['accuracy'])
model.summary()
history = model.fit(X_train, y_train,
batch_size=batch_size,
epochs=epochs,
verbose=1,
validation_data=(X_val, y_val))
score = model.evaluate(X_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
#get the predictions for the test data
predicted_classes = model.predict_classes(X_test)

```

```
#get the indices to be plotted
y_true = data_test.iloc[:, 0]
from sklearn.metrics import classification_report
target_names = ["Class {}".format(i) for i in range(num_classes)]
print(classification_report(y_true, predicted_classes, target_names=target_names))
```

## Epochs:

Epoch	Accuracy	Loss
5	0.9828000068664551	0.06191356107592583
10	0.991100013256073	0.03980569168925285
15	0.9883000254631042	0.0770006999373436
30	0.9897000193595886	0.09153448790311813

- We will choose number of epochs = 10 because it has high accuracy, take little time and good loss.

## Learning Rates:

Learning Rate	Accuracy	Loss
0.001	0.991100013256073	0.03980569168925285
0.005	0.980400025844574	0.10530269891023636
0.05	0.10279999673366547	2.3154354095458984

- We will choose learning rate = 0.001 because it has high accuracy, take little time and good loss.

# Models:

- **Model 1 :**

```
✓ batch_size = 32

num_classes = 10

epochs = 10

#input image dimensions

img_rows, img_cols = 28, 28


model = Sequential()

model.add(Conv2D(32, kernel_size=(3, 3),
activation='relu',
kernel_initializer='he_normal',
input_shape=input_shape))

model.add(MaxPooling2D(pool_size=(2, 2),strides=2))

model.add(Conv2D(64, (3, 3), activation='relu'))

model.add(MaxPooling2D(pool_size=(2, 2),strides=2))

model.add(Conv2D(128, (3, 3), activation='relu'))

model.add(Flatten())

model.add(Dense(128, activation='relu'))


model.add(Dense(num_classes, activation='softmax'))

opt = keras.optimizers.Adam(learning_rate=0.001)

model.compile(optimizer=opt, loss='categorical_crossentropy', m
etrics=['accuracy'])
```

- ✓ Test loss: 0. 03980569168925285
- ✓ Test accuracy: 0. 991100013256073
- ✓ Total params: 241,546
- ✓ Trainable params: 241,546
- ✓ Non-trainable params: 0

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 128)	73856
flatten (Flatten)	(None, 1152)	0
dense (Dense)	(None, 128)	147584
dense_1 (Dense)	(None, 10)	1290
Total params: 241,546		
Trainable params: 241,546		

## • Model 2 :

```

✓ batch_size = 32

num_classes = 10

epochs = 10

#input image dimensions
img_rows, img_cols = 28, 28

model = Sequential()

model.add(Conv2D(32, kernel_size=(3, 3),
activation='relu',
kernel_initializer='he_normal',
input_shape=input_shape))

model.add(MaxPooling2D(pool_size=(2, 2), strides=2))

model.add(Conv2D(64, (3, 3), activation='relu'))

model.add(MaxPooling2D(pool_size=(2, 2), strides=2))

model.add(Flatten())

model.add(Dense(64, activation='relu'))

model.add(Dense(num_classes, activation='softmax'))

opt = keras.optimizers.Adam(learning_rate=0.001)

```

```
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
```

- ✓ Test accuracy: 0.9886999726295471
- ✓ Test loss: 0.048302967101335526
- ✓ Total params: 121,930
- ✓ Trainable params: 121,930
- ✓ Non-trainable params: 0

Model: "sequential"		
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten (Flatten)	(None, 1600)	0
dense (Dense)	(None, 64)	102464
dense_1 (Dense)	(None, 10)	650
Total params: 121,930		
Trainable params: 121,930		

### • Model 3:

```
✓ batch_size = 32

num_classes = 10

epochs = 10

#input image dimensions

img_rows, img_cols = 28, 28

model = Sequential()

model.add(Conv2D(32, kernel_size=(3, 3),
activation='relu',
kernel_initializer='he_normal',
input_shape=input_shape))

model.add(MaxPooling2D(pool_size=(2, 2), strides=2))

model.add(Conv2D(128, (3, 3), activation='relu'))

model.add(Flatten())

model.add(Dense(128, activation='relu'))
```

```

model.add(Dense(num_classes, activation='softmax'))

opt = keras.optimizers.Adam(learning_rate=0.001)

model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])

```

- 
- ✓ Loss: 0.07243524491786957
- ✓ Accuracy : 0.989300012588501
- ✓ Total params: 2,021,194
- ✓ Trainable params: 2,021,194
- ✓ Non-trainable params: 0

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 128)	36992
flatten (Flatten)	(None, 15488)	0
dense (Dense)	(None, 128)	1982592
dense_1 (Dense)	(None, 10)	1290

## • Model 4:

```

✓ batch_size = 32

num_classes = 10

epochs = 10

#input image dimensions
img_rows, img_cols = 28, 28

model = Sequential()

model.add(Conv2D(30, kernel_size=(3, 3),
activation='relu',
kernel_initializer='he_normal',
input_shape=input_shape))

model.add(MaxPooling2D(pool_size=(2, 2), strides=2))

```



```

model.add(Flatten())

model.add(Dense(30, activation='relu'))

model.add(Dense(num_classes, activation='softmax'))

opt = keras.optimizers.Adam(learning_rate=0.001)

model.compile(optimizer=opt, loss='categorical_crossentropy', me
trics=['accuracy'])

```

- ✓ Loss: 0.051949601620435715
- ✓ Accuracy: 0.984499990940094
- ✓ Total params: 152,740
- ✓ Trainable params: 152,740
- ✓ Non-trainable params: 0

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 30)	300
max_pooling2d (MaxPooling2D)	(None, 13, 13, 30)	0
flatten (Flatten)	(None, 5070)	0
dense (Dense)	(None, 30)	152130
dense_1 (Dense)	(None, 10)	310
Total params: 152,740		

## • Model 5:

- ✓ `batch_size = 32`
- `num_classes = 10`
- `epochs = 10`
- `#input image dimensions`
- `img_rows, img_cols = 28, 28`
- `model = Sequential()`
- `model.add(Conv2D(64, kernel_size=(3, 3),`
- `activation='relu',`
- `kernel_initializer='he_normal',`

```

input_shape=input_shape))

model.add(MaxPooling2D(pool_size=(2, 2),strides=2))

model.add(Conv2D(128, (3, 3), activation='relu'))

model.add(MaxPooling2D(pool_size=(2, 2),strides=2))

model.add(Flatten())

model.add(Dense(128, activation='relu'))

model.add(Dense(num_classes, activation='softmax'))

opt = keras.optimizers.Adam(learning_rate=0.001)

model.compile(optimizer=opt, loss='categorical_crossentropy',
              metrics=['accuracy'])

```

- ✓ Loss: 0.0526624396443367
- ✓ Accuracy: 0.9891999959945679
- ✓ Total params: 485,514
- ✓ Trainable params: 485,514
- ✓ Non-trainable params: 0

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d (MaxPooling2D)	(None, 13, 13, 64)	0
conv2d_1 (Conv2D)	(None, 11, 11, 128)	73856
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 128)	0
flatten (Flatten)	(None, 3200)	0
dense (Dense)	(None, 128)	409728
dense_1 (Dense)	(None, 10)	1290

- We will choose Model 1 because it has high accuracy = %99.11 , take little time , good loss = %3.9 and minimum number of parameters .

## Batch Sizes:

Batch Size	Accuracy	Loss
32	0.991100013256073	0.03980569168925285
64	0.9897000193595886	0.04299784079194069
128	0.9908000230789185	0.03621158003807068

- We will choose batch size = 32 because it has high accuracy, take little time and good loss.

## Activation Functions:

Activation Function	Accuracy	Loss
ReLU	0.991100013256073	0.03980569168925285
tanh	0.9901000261306763	0.03669922053813934
sigmoid	0.9904999732971191	0.03228173032402992

- We will choose activation function: ReLU because it has high accuracy, take little time and good loss.

## Optimizers:

Optimizer	Accuracy	Loss
Adam	0.991100013256073	0.03980569168925285
SGD	0.9729999899864197	0.09458785504102707
Adamax	0.9914000034332275	0.030355028808116913

- We will choose optimizer: Adamax because it has high accuracy, take little time and good loss.

## Dropout:

- **First:**

```
batch_size = 32
num_classes = 10
epochs = 10

#input image dimensions
img_rows, img_cols = 28, 28

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
activation='relu',
kernel_initializer='he_normal',
input_shape=input_shape))

model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
    ➤ model.add(Dropout(0.25))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
    ➤ model.add(Dropout(0.4))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
    ➤ model.add(Dropout(0.3))
model.add(Dense(num_classes, activation='softmax'))
opt = keras.optimizers.Adamax(learning_rate=0.001)
```

```
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
```

- ✓ Loss: 0.021553359925746918
- ✓ Test accuracy: 0.9929999709129333
- ✓ When we add `model.add(Dropout(0.25))` eliminate about a quarter of the neurans from the first conv2D layer.
- ✓ When we add `model.add(Dropout(0.4))` eliminate about a %40 of the neurans from the second conv2D layer.
- ✓ When we add `model.add(Dropout(0.3))` eliminate about a %30 of the neurans from the third conv2D layer.

## • Second:

```
batch_size = 32
num_classes = 10
epochs = 10

#input image dimensions
img_rows, img_cols = 28, 28

model = Sequential()

model.add(Conv2D(32, kernel_size=(3, 3),
activation='relu',
kernel_initializer='he_normal',
input_shape=input_shape))

model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
    ➤ model.add(Dropout(0.35))

model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(Flatten())

model.add(Dense(128, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))

opt = keras.optimizers.Adamax(learning_rate=0.001)

model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
```

- ✓ Loss: 0.026227835565805435
- ✓ Test accuracy: 0.9915000200271606
- ✓ When we add `model.add(Dropout(0.35))` eliminate about a %35 of the neurans from the first conv2D layer.

### • Third:

```
batch_size = 32
num_classes = 10
epochs = 10

#input image dimensions
img_rows, img_cols = 28, 28

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
activation='relu',
kernel_initializer='he_normal',
input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
    ➤ model.add(Dropout(0.15))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(Flatten())
model.add(Dense(128, activation='relu'))

model.add(Dense(num_classes, activation='softmax'))
opt = keras.optimizers.Adamax(learning_rate=0.001)
model.compile(optimizer=opt, loss='categorical_crossentropy', metric
s=['accuracy'])
```

- ✓ Loss: 0.028334137052297592
- ✓ Test accuracy: 0.991599977016449

- ✓ When we add `model.add(Dropout(0.15))` eliminate about a %15 of the neurans from the second conv2D layer.

- **Fours:**

```
#m_1
batch_size = 32
num_classes = 10
epochs = 10
#input image dimensions
img_rows, img_cols = 28, 28

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
activation='relu',
kernel_initializer='he_normal',
input_shape=input_shape))

model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
    ➤ model.add(Dropout(0.35))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
    ➤ model.add(Dropout(0.70))
model.add(Dense(num_classes, activation='softmax'))
opt = keras.optimizers.Adamax(learning_rate=0.001)
model.compile(optimizer=opt, loss='categorical_crossentropy', metri
cs=['accuracy'])
```

- ✓ Loss: 0.03061206452548504
- ✓ Test accuracy: 0.9904000163078308
- ✓ When we add `model.add(Dropout(0.35))` eliminate about a %35 of the neurans from the first conv2D layer.

- ✓ When we add `model.add(Dropout(0.70))` eliminate about a %70 of the neurans from the third conv2D layer.

We will choose first dropout because it has high accuracy, take little time and good loss

- **Best model with test batch size:**

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
from subprocess import check_output
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
import matplotlib.pyplot as plt
import matplotlib
import matplotlib.image as mpimg
import numpy as np
from numpy import random
import keras
from keras.utils import to_categorical
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten, Conv2D, MaxPooling2D
from keras.utils import np_utils
from keras.preprocessing.image import ImageDataGenerator
from keras.callbacks import ReduceLROnPlateau, Callback
from keras import regularizers
from keras.optimizers import Adam
import tensorflow as tf
from keras.optimizers import SGD
from sklearn.utils import shuffle

data_train = pd.read_csv('/content/mnist_train.csv')
data_test = pd.read_csv('/content/mnist_test.csv')
print("data_train" , data_train.shape)
print("data_test" , data_test.shape)
img_rows, img_cols = 28, 28
input_shape = (img_rows, img_cols, 1)
X = np.array(data_train.iloc[:, 1:])
y = to_categorical(np.array(data_train.iloc[:, 0]))
#Here we split validation data to optimize classifier during training
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=13)
#Test data
X_test = np.array(data_test.iloc[:, 1:])
```



```

y_test = to_categorical(np.array(data_test.iloc[:, 0]))
X_train = np.array(X_train)
y_train = np.array(y_train)
X_test = np.array(X_test)
y_test = np.array(y_test)
X_train, y_train = shuffle(X_train, y_train)
X_test, y_test = shuffle(X_test, y_test)
X_train = X_train.reshape(X_train.shape[0], img_rows, img_cols, 1)
X_test = X_test.reshape(X_test.shape[0], img_rows, img_cols, 1)
X_val = X_val.reshape(X_val.shape[0], img_rows, img_cols, 1)
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_val = X_val.astype('float32')
X_train /= 255
X_test /= 255
X_val /= 255

print(X_train)
print("X_train:{}\ny_train:{}\n\nX_val:{}\ny_val:{}\n\nX_test:{}\ny_
_test:{}".format
(X_train.shape, y_train.shape, X_val.shape, y_val.shape, X_test.sha
pe, y_test.shape))
#m_1

batch_size = 32

num_classes = 10

epochs = 10

#input image dimensions

img_rows, img_cols = 28, 28

model = Sequential()

model.add(Conv2D(32, kernel_size=(3, 3),
activation='relu',
kernel_initializer='he_normal',
input_shape=input_shape))

model.add(MaxPooling2D(pool_size=(2, 2),strides=2))

model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3), activation='relu'))

model.add(MaxPooling2D(pool_size=(2, 2),strides=2))

model.add(Dropout(0.4))

```

```

model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(num_classes, activation='softmax'))
opt = keras.optimizers.Adamax(learning_rate=0.001)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])

```

```

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
-----		
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
-----		
dropout (Dropout)	(None, 13, 13, 32)	0
-----		
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
-----		
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
-----		
dropout_1 (Dropout)	(None, 5, 5, 64)	0
-----		
conv2d_2 (Conv2D)	(None, 3, 3, 128)	73856
-----		
flatten (Flatten)	(None, 1152)	0
-----		
dense (Dense)	(None, 128)	147584
-----		
dropout_2 (Dropout)	(None, 128)	0
-----		
dense_1 (Dense)	(None, 10)	1290
=====		

Total params: 241,546  
Trainable params: 241,546  
Non-trainable params: 0

```
history = model.fit(X_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(X_val, y_val))
score = model.evaluate(X_test, y_test, verbose=0 , batch_size=128)
```

Epoch 1/10

1500/1500 [=====] - 54s 36ms/step - loss: 0.6774 - accuracy: 0.7749 -  
val\_loss: 0.0955 - val\_accuracy: 0.9693

Epoch 2/10

1500/1500 [=====] - 53s 36ms/step - loss: 0.1390 - accuracy: 0.9576 -  
val\_loss: 0.0616 - val\_accuracy: 0.9806

Epoch 3/10

1500/1500 [=====] - 53s 35ms/step - loss: 0.0861 - accuracy: 0.9742 -  
val\_loss: 0.0488 - val\_accuracy: 0.9852

Epoch 4/10

1500/1500 [=====] - 53s 36ms/step - loss: 0.0739 - accuracy: 0.9763 -  
val\_loss: 0.0453 - val\_accuracy: 0.9850

Epoch 5/10

1500/1500 [=====] - 53s 36ms/step - loss: 0.0665 - accuracy: 0.9792 -  
val\_loss: 0.0374 - val\_accuracy: 0.9881

Epoch 6/10

1500/1500 [=====] - 53s 35ms/step - loss: 0.0545 - accuracy: 0.9829 -  
val\_loss: 0.0357 - val\_accuracy: 0.9887

Epoch 7/10

1500/1500 [=====] - 53s 35ms/step - loss: 0.0525 - accuracy: 0.9839 -  
val\_loss: 0.0373 - val\_accuracy: 0.9883

Epoch 8/10

1500/1500 [=====] - 53s 35ms/step - loss: 0.0457 - accuracy: 0.9858 -  
val\_loss: 0.0339 - val\_accuracy: 0.9891

Epoch 9/10

1500/1500 [=====] - 53s 35ms/step - loss: 0.0433 - accuracy: 0.9867 -  
val\_loss: 0.0314 - val\_accuracy: 0.9902

Epoch 10/10

1500/1500 [=====] - 52s 34ms/step - loss: 0.0384 - accuracy: 0.9877 -  
val\_loss: 0.0314 - val\_accuracy: 0.9908

```
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Test loss: 0.023057807236909866

Test accuracy: 0.9926000237464905

```
#get the predictions for the test data
predicted_classes = model.predict_classes(X_test)
#get the indices to be plotted
y_true = data_test.iloc[:, 0]
from sklearn.metrics import classification_report
target_names = ["Class {}".format(i) for i in range(num_classes)]
print(classification_report(y_true, predicted_classes, target_names
= target_names))
```

	precision	recall	f1-score	support
Class 0	0.09	0.09	0.09	980
Class 1	0.13	0.13	0.13	1135
Class 2	0.12	0.12	0.12	1032
Class 3	0.10	0.10	0.10	1010
Class 4	0.09	0.09	0.09	982
Class 5	0.09	0.09	0.09	892
Class 6	0.10	0.10	0.10	958
Class 7	0.10	0.10	0.10	1028
Class 8	0.10	0.10	0.10	974
Class 9	0.11	0.11	0.11	1009
accuracy			0.10	10000
macro avg	0.10	0.10	0.10	10000
weighted avg	0.10	0.10	0.10	10000

## Summary:

Epochs	Learning Rate	Batch size	Activation function	optimizers	Accuracy	Loss
5	0.001	32	relu	Adam	0.982	0.061
10	0.001	32	relu	Adam	0.991	0.039
15	0.001	32	relu	Adam	0.988	0.077

<b>30</b>	<b>0.001</b>	<b>32</b>	<b>relu</b>	<b>Adam</b>	<b>0.989</b>	<b>0.091</b>
<b>10</b>	<b>0.005</b>	<b>32</b>	<b>relu</b>	<b>Adam</b>	<b>0.980</b>	<b>0.105</b>
<b>10</b>	<b>0.05</b>	<b>32</b>	<b>relu</b>	<b>Adam</b>	<b>0.102</b>	<b>2.315</b>
<b>10</b>	<b>0.001</b>	<b>64</b>	<b>relu</b>	<b>Adam</b>	<b>0.989</b>	<b>0.042</b>
<b>10</b>	<b>0.001</b>	<b>128</b>	<b>relu</b>	<b>Adam</b>	<b>0.990</b>	<b>0.036</b>
<b>10</b>	<b>0.001</b>	<b>32</b>	<b>tanh</b>	<b>Adam</b>	<b>0.990</b>	<b>0.036</b>
<b>10</b>	<b>0.001</b>	<b>32</b>	<b>sigmoid</b>	<b>Adam</b>	<b>0.990</b>	<b>0.032</b>
<b>10</b>	<b>0.001</b>	<b>32</b>	<b>relu</b>	<b>SGD</b>	<b>0.972</b>	<b>0.0945</b>
<b>10</b>	<b>0.001</b>	<b>32</b>	<b>relu</b>	<b>Adamax</b>	<b>0.9914</b>	<b>0.0303</b>